

Addison J. Wu

addisonwu@princeton.edu

addisonwu05.github.io

Citizenship: USA, Canada

Education

2023 – 2027 **Princeton University**

BSE in Computer Science, Minor in Mathematics – GPA: 3.98/4.00

Honors: Exemplary Independent Work Award, Shapiro Prize for Academic Excellence, NeurIPS Scholar Award, Tau Beta Pi, Sigma Xi (Associate)

Coursework: addisonwu05.github.io/academics

Experience

2026 – [Center for AI Safety](#)

Research Intern, advised by Mantas Mazeika, Dan Hendrycks

2024 – [Princeton Language and Intelligence](#)

Undergraduate Researcher, advised by Tom Griffiths, Zhuang Liu

Interactive evaluations for model behavior; data and training constraints

2023 – 2025 [Multi-Physics Interaction Lab](#), University of Waterloo

Undergraduate Researcher, advised by Jean-Pierre Hickey

Neural networks for atmospheric infrasound localization

Publications

 [Google Scholar](#) * → Equal contribution

Conference/Journal Papers

1. **Large Language Models Develop Novel Social Biases Through Adaptive Exploration**
Addison J. Wu*, Ryan Liu*, Xuechunzi Bai, Thomas L. Griffiths
ICML 2026 **Oral** [\[link\]](#)
2. **Ads in AI Chatbots? An Analysis of How Large Language Models Navigate Conflicts of Interest**
Addison J. Wu*, Ryan Liu*, Shuyue Stella Li, Yulia Tsvetkov, Thomas L. Griffiths
COLM 2026 [\[link\]](#)
3. **Query Timing Produces Opposite Positional Biases Between LLMs and Humans**
Jasin Cekinmez*, Addison J. Wu*, Thomas L. Griffiths
EMNLP 2026 Findings. ICBINB @ ICLR 2026 Entropic Paper Award [\[link\]](#)
4. **Are Large Language Models Sensitive to the Motives Behind Communication?**
Addison J. Wu*, Ryan Liu*, Kerem Oktar*, Theodore R. Sumers, Thomas L. Griffiths
NeurIPS 2025. PragLM @ COLM 2025 Oral [\[link\]](#)
5. **Mind Your Step (by Step): Chain-of-Thought can Reduce Performance on Tasks where Thinking Makes Humans Worse**
Ryan Liu*, Jiayi Geng*, Addison J. Wu, Ilia Sucholutsky, Tania Lombrozo, Thomas L. Griffiths
ICML 2025 [\[link\]](#)
6. **Toward a Neural Network-Based Approach for Improved Atmospheric Infrasound Localization**
Addison J. Wu, Arnav Joshi, Jean-Pierre Hickey
IEEE Access, 2025 [\[link\]](#)

Workshop Papers

1. **Where did the ambiguity go? Examining how multimodal models interpret polysemous words**
Jasin Cekinmez*, **Addison J. Wu***, Raja Marjeh, Thomas L. Griffiths
SciFM @ COLM 2026 *Oral* [\[link\]](#)
2. **Quantifying Empirical Compute-Supervision Tradeoffs in RLVR**
Ryo Mitsuhashi*, Patrick Chen*, Isabelle Tseng*, Jasin Cekinmez*, **Addison J. Wu***
Combining Theory and Benchmarks @ ICML 2026 [\[link\]](#)
3. **OccludeBench: Can VLMs See the Hint?**
Jasin Cekinmez*, **Addison J. Wu***, Ryo Mitsuhashi*, Linrong Cai, Xingyu Fu, Zhuang Liu
Computer Vision in the Wild @ CVPR 2026 [\[link\]](#)

Preprints

1. **SportD: Can VLMs Physically Strategize?**
Jasin Cekinmez*, **Addison J. Wu***, Haotian Xia, Akshaya Bharadhwaj, Anay Putty, Anirudh Ravishankar, Jaewoong Lee, Jinglin Xiao, Kyumin Andrew Shim, Mishika Ahuja, Nisarga Patil, Leo Liu, Zhuohan Liu, Weining Shen
arXiv preprint, 2026 [\[link\]](#)
2. **Guess the Unified Model: How Much Can We Recover from Generated Images?**
Jasin Cekinmez*, Ryo Mitsuhashi*, **Addison J. Wu***, Yida Yin
arXiv preprint, 2026 [\[link\]](#)

Policy Citations

- 2026 U.S. Senate Judiciary Committee, Subcommittee on Crime and Counterterrorism
Ads in AI Chatbots? cited in the written testimony of Lindsay Owens, Groundwork Collaborative, *Your Data, Their Profit: The Consumer Cost of AI Surveillance Pricing*

Press Coverage

- 2026 The Guardian, [Will AI give you the job? Automated hiring tools spark discrimination and secrecy lawsuits](#)
- 2026 Forbes, [AI Hiring Tools Don't Just Learn Bias, AI Forms New Biases Of Its Own](#)
- 2026 MIT Technology Review, [AI is more likely than humans to form biases when hiring](#)
- 2026 The Times of India, [AI can become more biased than humans while choosing who gets hired; study finds](#)
- 2026 Gizmodo, [AI Tends to Develop New Stereotypes to Base Hiring Decisions On, Study Says](#)
- 2026 Arizona's Family, [AI can invent its own hiring bias — and it happens fast, new study finds](#)
- 2026 Indian Express, [Your AI 'assistant' is making money for someone else, you are the one paying. Here's exactly how](#)
- 2026 Hacks/Hackers, [New research: 18 of 23 AI models prioritize company revenue over users when ads enter the picture](#)
- 2025 The Information, [Where LLMs Are Falling Short](#) [\[full text\]](#)

Invited and Contributed Talks

- 2026 **ICML 2026 Oral**
Large Language Models Develop Novel Social Biases Through Adaptive Exploration
- 2026 **SPSP 2026 Invited Symposia**
Large Language Models Develop Novel Social Biases Through Adaptive Exploration
(presented by co-author X. Bai)

- 2025 **UCLA NLP Seminar Series**
Are Large Language Models Sensitive to the Motives Behind Communication?
Large Language Models Develop Novel Social Biases Through Adaptive Exploration
(presented by co-first author R. Liu)
- 2025 **PragLM @ COLM 2025**
Are Large Language Models Sensitive to the Motives Behind Communication?

Teaching Experience

- 2024–2025 Undergraduate Course Assistant, Algorithms and Data Structures (COS 226)
Fall 2024, Spring 2025, Fall 2025

Professional Services

- 2026 Reviewer, NeurIPS 2026
2025 Reviewer, MTI-LLM @ NeurIPS 2025