

# Addison J. Wu

[addisonwu@princeton.edu](mailto:addisonwu@princeton.edu)

[addisonwu05.github.io](https://addisonwu05.github.io)

Citizenship: USA, Canada

## Education

2023 – 2027 **Princeton University**  
*BSE in Computer Science, Minor in Mathematics – GPA: 3.98/4.00*  
Coursework: [addisonwu05.github.io/academics](https://addisonwu05.github.io/academics)

## Experience

2026 – [Center for AI Safety](#)  
*Research Intern*, advised by Mantas Mazeika, Dan Hendrycks

2024 – [Princeton Language and Intelligence](#)  
*Undergraduate Researcher*, advised by Tom Griffiths, Zhuang Liu  
Interactive evaluations for model behavior; data and training constraints

2023 – 2025 [Multi-Physics Interaction Lab](#), University of Waterloo  
*Undergraduate Researcher*, advised by Jean-Pierre Hickey  
Neural networks for atmospheric infrasound localization

## Honors

- Oral Presentation, 2026 (ICML)
- Exemplary Independent Work Award, 2026 (Princeton University)
- Entropic Paper Award, 2026 (ICBINB @ ICLR)
- NeurIPS Scholar Award, 2025
- OURSIP Research Award, funded by the Hewlett Foundation, 2025 (Princeton University)
- Tau Beta Pi engineering honor society, 2025 (Princeton University)
- Sigma Xi research honor society, associate member, 2025 (Princeton University)
- Shapiro Prize for Academic Excellence, 2024 (Princeton University)

## Publications

 [Google Scholar](#) \* → Equal contribution

### Conference/Journal Papers

1. **Large Language Models Develop Novel Social Biases Through Adaptive Exploration**  
**Addison J. Wu\***, Ryan Liu\*, Xuechunzi Bai, Thomas L. Griffiths  
*ICML 2026 Oral* [\[link\]](#)
2. **Ads in AI Chatbots? An Analysis of How Large Language Models Navigate Conflicts of Interest**  
**Addison J. Wu\***, Ryan Liu\*, Shuyue Stella Li, Yulia Tsvetkov, Thomas L. Griffiths  
*COLM 2026* [\[link\]](#)
3. **Query Timing Produces Opposite Positional Biases Between LLMs and Humans**  
Jasin Cekinmez\*, **Addison J. Wu\***, Thomas L. Griffiths  
*EMNLP 2026 Findings; ICBINB @ ICLR 2026 Entropic Paper Award* [\[link\]](#)
4. **Are Large Language Models Sensitive to the Motives Behind Communication?**  
**Addison J. Wu\***, Ryan Liu\*, Kerem Oktar\*, Theodore R. Sumers, Thomas L. Griffiths  
*NeurIPS 2025. PragLM @ COLM 2025 Oral* [\[link\]](#)
5. **Mind Your Step (by Step): Chain-of-Thought can Reduce Performance on Tasks where Thinking Makes Humans Worse**  
Ryan Liu\*, Jiayi Geng\*, **Addison J. Wu**, Ilia Sucholutsky, Tania Lombrozo, Thomas L. Griffiths  
*ICML 2025* [\[link\]](#)
6. **Toward a Neural Network-Based Approach for Improved Atmospheric Infrasound Localization**  
**Addison J. Wu**, Arnav Joshi, Jean-Pierre Hickey  
*IEEE Access*, 2025 [\[link\]](#)

## Workshop Papers

1. **Where did the ambiguity go? Examining how multimodal models interpret polysemous words**  
Jasin Cekinmez\*, **Addison J. Wu\***, Raja Marjeh, Thomas L. Griffiths  
*SciFM @ COLM 2026 Oral* [\[link\]](#)
2. **Quantifying Empirical Compute-Supervision Tradeoffs in RLVR**  
Ryo Mitsuhashi\*, Patrick Chen\*, Isabelle Tseng\*, Jasin Cekinmez\*, **Addison J. Wu\***  
*Combining Theory and Benchmarks @ ICML 2026* [\[link\]](#)
3. **OccludeBench: Can VLMs See the Hint?**  
Jasin Cekinmez\*, **Addison J. Wu\***, Ryo Mitsuhashi\*, Linrong Cai, Xingyu Fu, Zhuang Liu  
*Computer Vision in the Wild @ CVPR 2026* [\[link\]](#)

## Preprints

1. **SportD: Can VLMs Physically Strategize?**  
Jasin Cekinmez\*, **Addison J. Wu\***, Haotian Xia, Akshaya Bharadhwaj, Anay Putty, Anirudh Ravishankar, Jaewoong Lee, Jinglin Xiao, Kyumin Andrew Shim, Mishika Ahuja, Nisarga Patil, Leo Liu, Zhuohan Liu, Weining Shen  
arXiv preprint, 2026 [\[link\]](#)
2. **Guess the Unified Model: How Much Can We Recover from Generated Images?**  
Jasin Cekinmez\*, Ryo Mitsuhashi\*, **Addison J. Wu\***, Yida Yin  
arXiv preprint, 2026 [\[link\]](#)

## Policy Citations

- 2026                    U.S. Senate Judiciary Committee, Subcommittee on Crime and Counterterrorism  
*Ads in AI Chatbots?* cited in the written testimony of Lindsay Owens, Groundwork Collaborative, [Your Data, Their Profit: The Consumer Cost of AI Surveillance Pricing](#)

## Media Coverage

- 2026                    The Guardian, [Will AI give you the job? Automated hiring tools spark discrimination and secrecy lawsuits](#)
- 2026                    Forbes, [AI Hiring Tools Don't Just Learn Bias, AI Forms New Biases Of Its Own](#)
- 2026                    MIT Technology Review, [AI is more likely than humans to form biases when hiring](#)
- 2026                    The Times of India, [AI can become more biased than humans while choosing who gets hired; study finds](#)
- 2026                    Gizmodo, [AI Tends to Develop New Stereotypes to Base Hiring Decisions On, Study Says](#)
- 2026                    Arizona's Family, [AI can invent its own hiring bias — and it happens fast, new study finds](#)
- 2026                    Indian Express, [Your AI 'assistant' is making money for someone else, you are the one paying. Here's exactly how](#)
- 2026                    Hacks/Hackers, [New research: 18 of 23 AI models prioritize company revenue over users when ads enter the picture](#)
- 2025                    The Information, [Where LLMs Are Falling Short](#) [\[full text\]](#)

## Teaching Experience

- 2024–2025            Undergraduate Course Assistant, Algorithms and Data Structures (COS 226)  
*Fall 2024, Spring 2025, Fall 2025*

## Professional Services

- 2026                    Reviewer, NeurIPS 2026
- 2025                    Reviewer, MTI-LLM @ NeurIPS 2025