Addison J. Wu

addisonwu@princeton.edu addisonwu05.github.io 609-937-4841 Citizenship: USA, Canada

Education

2023 – 2027 **Princeton University**
BSE in Computer Science, Minor in Mathematics – GPA: 3.98/4.00
Coursework: addisonwu05.github.io/academics

Experience

2024 – Princeton Language and Intelligence
Undergraduate Researcher, advised by Tom Griffiths, Zhuang Liu
Interactive evaluations for model behavior; data and training constraints

Jun–Aug 2026 Center for AI Safety, San Francisco
Research Intern, advised by Mantas Mazeika, Dan Hendrycks

2023 – 2025 Multi-Physics Interaction Lab, University of Waterloo
Undergraduate Researcher, advised by Jean-Pierre Hickey
Neural networks for atmospheric infrasound localization

Honors

- Oral Presentation, 2026 (ICML)
- Exemplary Independent Work Award, 2026 (Princeton University)
- Entropic Paper Award, 2026 (ICBINB @ ICLR)
- NeurIPS Scholar Award, 2025 (NeurIPS)
- OURSIP Research Award, funded by the Hewlett Foundation, 2025 (Princeton University)
- Tau Beta Pi engineering honor society, 2025 (Princeton University)
- Sigma Xi research honor society, associate member, 2025 (Princeton University)
- Shapiro Prize for Academic Excellence, 2024 (Princeton University)

Publications

Google Scholar $* \rightarrow$ Equal contribution

Conference/Journal Papers

1. **Large Language Models Develop Novel Social Biases Through Adaptive Exploration**
Addison J. Wu*, Ryan Liu*, Xuechunzi Bai, Thomas L. Griffiths
*ICML 2026 **Oral*** [link]
2. **Ads in AI Chatbots? An Analysis of How Large Language Models Navigate Conflicts of Interest**
Addison J. Wu*, Ryan Liu*, Shuyue Stella Li, Yulia Tsvetkov, Thomas L. Griffiths
COLM 2026 [link]
3. **Query Timing Produces Opposite Positional Biases Between LLMs and Humans**
Jasin Cekinmez*, **Addison J. Wu***, Thomas L. Griffiths
*EMNLP 2026 Findings; ICBINB @ ICLR 2026 **Entropic Paper Award*** [link]
4. **Are Large Language Models Sensitive to the Motives Behind Communication?**
Addison J. Wu*, Ryan Liu*, Kerem Oktar*, Theodore R. Sumers, Thomas L. Griffiths
*NeurIPS 2025; PragLM @ COLM 2025 **Oral*** [link]
5. **Mind Your Step (by Step): Chain-of-Thought can Reduce Performance on Tasks where Thinking Makes Humans Worse**
Ryan Liu*, Jiayi Geng*, **Addison J. Wu**, Ilia Sucholutsky, Tania Lombrozo, Thomas L. Griffiths
ICML 2025 [link]
6. **Toward a Neural Network-Based Approach for Improved Atmospheric Infrasound Localization**
Addison J. Wu, Arnav Joshi, Jean-Pierre Hickey
IEEE Access, 2025 [link]

Workshop Papers

1. **Where did the ambiguity go? Examining how multimodal models interpret polysemous words**
   Jasin Cekinmez*, **Addison J. Wu***, Raja Marjieh, Thomas L. Griffiths
   SciFM @ COLM 2026 ***Oral*** [link]
2. **Quantifying Empirical Compute-Supervision Tradeoffs in RLVR**
   Ryo Mitsuhashi*, Patrick Chen*, Isabelle Tseng*, Jasin Cekinmez*, **Addison J. Wu***
   Combining Theory and Benchmarks @ ICML 2026 [link]
3. **OccludeBench: Can VLMs See the Hint?**
   Jasin Cekinmez*, **Addison J. Wu***, Ryo Mitsuhashi*, Linrong Cai, Xingyu Fu, Zhuang Liu
   Computer Vision in the Wild @ CVPR 2026 [link]

Preprints

1. **SportD: Can VLMs Physically Strategize?**
   Jasin Cekinmez*, **Addison J. Wu***, Haotian Xia, Akshaya Bharadhwaj, Anay Putty, Anirudh Ravishankar, Jaewoong Lee, Jinglin Xiao, Kyumin Andrew Shim, Mishika Ahuja, Nisarga Patil, Leo Liu, Zhuohan Liu, Weining Shen
   arXiv preprint, 2026 [link]
2. **Guess the Unified Model: How Much Can We Recover from Generated Images?**
   Jasin Cekinmez*, Ryo Mitsuhashi*, **Addison J. Wu***, Yida Yin
   arXiv preprint, 2026 [link]

Policy Citations

2026	U.S. Senate Judiciary Committee, Subcommittee on Crime and Counterterrorism *Ads in AI Chatbots?* cited in the written testimony of Lindsay Owens, Groundwork Collaborative, *Your Data, Their Profit: The Consumer Cost of AI Surveillance Pricing*

Media Coverage

2026	The Guardian, Will AI give you the job? Automated hiring tools spark discrimination and secrecy lawsuits
2026	Forbes, AI Hiring Tools Don't Just Learn Bias, AI Forms New Biases Of Its Own
2026	MIT Technology Review, AI is more likely than humans to form biases when hiring
2026	The Times of India, AI can become more biased than humans while choosing who gets hired; study finds
2026	Gizmodo, AI Tends to Develop New Stereotypes to Base Hiring Decisions On, Study Says
2026	Arizona's Family, AI can invent its own hiring bias — and it happens fast, new study finds
2026	Indian Express, Your AI 'assistant' is making money for someone else, you are the one paying. Here's exactly how
2026	Hacks/Hackers, New research: 18 of 23 AI models prioritize company revenue over users when ads enter the picture
2025	The Information, Where LLMs Are Falling Short [full text]

Teaching Experience

2024–2025	Undergraduate Course Assistant, Algorithms and Data Structures (COS 226) *Fall 2024, Spring 2025, Fall 2025*

Professional Services

2026	Reviewer, NeurIPS 2026
2025	Reviewer, MTI-LLM @ NeurIPS 2025